

28 September 2026

The Editors JMIR Mental Health

Dear Editor,

We are pleased to submit our original research manuscript, **“Feature Vectors and Narrative Embeddings for Predicting a Treatment Switching Proxy for Treatment Resistant Depression: Retrospective Cohort Study,”** for consideration in *JMIR Mental Health*.

The study asks whether pretrained embeddings of rule-based patient narratives predict a treatment-switching proxy for treatment-resistant depression (TRD) better than structured feature vectors built from the same electronic health records (EHRs). Prediction is made at the index antidepressant prescription.

In 42,579 patients from one community health system, the two approaches performed similarly. Embedded logistic regression reached a receiver operating characteristic area under the curve (ROC AUC) of 0.657 (95% CI 0.643-0.672). Feature-vector XGBoost reached 0.649 (95% CI 0.634-0.664). Their paired difference was 0.008 (95% CI -0.003 to 0.019). The classifier changed the result: embeddings improved logistic regression and reduced discrimination for all three tree ensembles. Logistic regression led the embedded models on all four encoders.

The predictive signal came from familiar clinical history. Permuting psychiatric history caused the largest loss, and medication burden and prior treatment also contributed. Permuting race/ethnicity or social determinants changed ROC AUC by no more than 0.003. Nearest-neighbor retrieval over the embedding showed that proximity tracks outcome risk, but it stayed below both leading classifiers at every neighborhood size. Its best ROC AUC was 0.625 (95% CI 0.610-0.641).

We are explicit that the outcome is a treatment-switching proxy rather than confirmed treatment resistance, that validation is internal, and that the subgroup analyses do not establish equitable performance. We think the results give a practical structured-feature benchmark for EHR embedding pipelines, and a clear case for validating the outcome before deploying either approach.

This work is reported in accordance with the **TRIPOD+AI** guideline for the reporting of clinical prediction models that use regression or machine learning methods; a completed TRIPOD+AI checklist is included with this submission.

We confirm that this manuscript describes original work, is not under consideration for publication elsewhere, and that all listed authors have read and approved the submission. The EHR data

are not publicly shareable; the full analysis code is available at <https://github.com/Pirate-Hunter-Zoro/TRD-EHR>. Funding, conflict of interest, and ethics/IRB statements are provided in the manuscript's Declarations: no conflicts of interest are declared; the work was funded by the William K. Warren Foundation, with no other funding source; and the study is a secondary analysis of de-identified data and is therefore not human-subjects research subject to institutional review board approval. All analysis was performed on local institutional hardware with a de-identified extract containing no protected health information, with no patient data transmitted to any external or commercial service.

Thank you for considering our work.

Sincerely,

Mikey Ferguson, BS Laureate Institute for Brain Research 6655 South Yale Avenue, Tulsa, OK 74136, United States mferguson@laureateinstitute.org on behalf of all authors